

# Hw1

## Task1

### 1.1 How to rasterize triangles in your own words.

#### 1.1.1 Compute the axis-aligned bounding box (AABB).

I take min/max of the three vertices in screen space, then shift by **-0.5 px** so that pixel indices represent **pixel centers**. Finally, I clamp the range to the framebuffer.

#### 1.1.2 Edge-Function Pre-computation

Each directed edge  $e_i = (v_i, v_j)$  is represented by a linear function

$$e(x, y) = ax + by + c$$

with

$$a = y_i - y_j, \quad b = x_j - x_i, \quad c = x_i y_j - x_j y_i.$$

#### 1.1.3 Consistent Orientation

The signed doubled area


$$A_2 = (x_1 - x_0)(y_2 - y_0) - (y_1 - y_0)(x_2 - x_0)$$

#### 1.1.4 Scanning the Bounding Box

The code iterates row by row through every pixel whose centre lies inside the bounding box. For each  $(p_x, p_y)$ , the centre coordinates are  $(x_s, y_s) = (p_x + 0.5, p_y + 0.5)$ . Computing  $e_0, e_1, e_2$  requires three multiply-add operations per edge—nine FLOPs total—followed by three comparisons. If the point passes the edge test, `fill_pixel` writes the colour to every subsample of that pixel.

## 1.2 Why the Algorithm Is No Worse than a Full AABB Scan

- I iterate over **exactly** the pixels inside the AABB, which is the same set examined by a naïve “sample-every-point-in-bbox” approach.
- Each pixel entails only **three multiply-adds and three comparisons**. No costly operations (div, sqrt) are used.
- Therefore the complexity is **O(bbox pixels)** in time and **O(1)** in extra memory.



## About Extra Credit

I tried to complete the extra credit for task 1, but ultimately was not able to finish it. To measure the time taken, I added the following code at line 307 in `src/drawrend.cpp`. However, I found that the time required to render the same image was unstable with each run. Can this kind of timing be used for comparison?

```
auto t0 = std::chrono::high_resolution_clock::now();
svg.draw(software_rasterizer, ndc_to_screen * svg_to_ndc[current_svg]);
auto t1 = std::chrono::high_resolution_clock::now();
double ms = std::chrono::duration<double, std::milli>(t1 - t0).count();
std::cout << ms << " ms\n";
```

## Task2

### 2.1 Supersampling algorithm & data structures

#### 2.1.1. sample\_rate

This integer controls how many sub-samples live inside each pixel—1, 4, 9, 16, and so on. You can change it at run time via `set_sample_rate()`, letting you trade rendering speed for image quality. Every buffer size allocation and every loop that iterates over sub-samples is driven by this value.

#### 2.1.2. sample\_buffer

`sample_buffer` is the central data structure that stores the intermediate colors of all sub-samples. It is a `vector<Color>` whose length equals  $\text{width} \times \text{height} \times \text{sample\_rate}$ , reserving `sample_rate` color slots for every pixel on the screen. The index scheme is simple: locate the pixel first, then offset by the sub-sample index— $\text{pixel\_id} * \text{sample\_rate} + \text{sub\_id}$ . During rasterization, we write directly into this buffer; we do **not** touch the 8-bit RGB framebuffer until the very end.

#### 2.1.3 Pipeline for a triangle

- a) **Bounding box clipping.**
- b) **Per-sub-sample test.**
  1. Iterate over all sub-samples inside each pixel.
  2. Evaluate edge functions at the floating-point positions.
  3. If inside, write the color directly to `sample_buffer`.
- c) **Resolve.** Average the  $n$  samples back to the 8-bit RGB framebuffer in `resolve_to_framebuffer()`.

## 2.2 Why supersampling is useful

- **Aliasing:** Single-point sampling cannot estimate partial coverage  $\rightarrow$  “jaggies”.
- **Coverage approximation:** With  $\geq 4$  samples, we can approximate the true area coverage by  $k/n$ .
- **Renderer-agnostic:** Pure geometry approach, independent of texture filters.

## 2.3 Modifications to the rasterization pipeline

Location	Change
<code>fill_pixel</code>	For points/lines, write the same color for every sub-sample.
<code>rasterize_triangle</code>	Added nested loops over sub-samples; kept fast path when $n=1$ .
<code>set_sample_rate</code> & <code>set_framebuffer_target</code>	Re-allocate <code>sample_buffer</code> .
<code>resolve_to_framebuffer</code>	Average $n$ sub-samples and write back to the RGB framebuffer.

## 2.4. Using supersampling to antialias triangles


For each pixel we count hit samples  $k$ .

- Completely inside:  $k=n$ .
- Completely outside:  $k=0$ .
- **Edge pixels:**  $0 < k < n$ . Final  $color = triangle\_color \times k/n + background \times (n - k)/n$ .  
When  $n$  increases  $1 \rightarrow 4 \rightarrow 16$ , the coverage levels become finer (25 %, 50 %, 75 %, ...) and the staircase edge progressively turns into a smooth gradient.

## 2.5. Results & explanation

The pixel inspector is placed over the skinny tip of the blue triangle:

- **1×** Only two colors (blue/white) → harsh jaggies.
- **4×** Mixed shades appear; 25/50/75 % coverage soften the edge.
- **16×** Even finer coverage steps; edge looks virtually smooth at normal viewing distance. Frame-rate drops because the inner loops run 16× more often.




### Additional Notes:

I completed the entire assignment under WSL. When I open the file with the default viewing parameters using


```
./draw ../svg/basic/test4.svg
```

I see a single blue pixel spilling outside the black frame. However, after I zoom in or out with the mouse wheel, all blue pixels are confined strictly within the frame.



## Task3

This picture looks like a robot dancing, though it could also be seen as throwing a softball.



### For extra Credit

#### Store the Rotation Angle

Add a float `view_rotation` field to the `DrawRend` class to store the current viewport rotation angle (in degrees). Initialize it to 0 in the constructor, `init()`, and `view_init()`.

#### Listen for Keyboard Events

In the switch statement inside `keyboard_event`, add:

```
case 'E': view_rotation -= 5.f; redraw(); break; // clockwise
case 'Q': view_rotation += 5.f; redraw(); break; // counterclockwise
```

Each key press updates the rotation angle and triggers a redraw.

#### Construct the Rotation Matrix

At the beginning of `redraw()`, create:

```
Matrix3x3 R = translate(0.5, 0.5) *
               rotate(view_rotation) *
               translate(-0.5, -0.5);
```

This applies a rotation around the center of the NDC (Normalized Device Coordinates) space, which is at (0.5, 0.5).

#### Insert Into the Matrix Stack


Replace the original drawing transformation, `ndc_to_screen * svg_to_ndc` with:

```
Matrix3x3 total = ndc_to_screen * R * svg_to_ndc[current_svg];
```

Apply this total matrix to all vertices (both the SVG content and the border lines), so the entire canvas rotates as the angle changes.

### Reset on View Initialization

In `view_init()`, reset `view_rotation` to 0 to ensure that pressing the space bar not only resets the view but also clears any rotation.





## Task 4

Inside a triangle, any point can be thought of as a blend of the three vertices in certain “proportions”; these three proportions are the point’s **barycentric coordinates**, denoted  $(w_0, w_1, w_2)$ .

As shown in the figure below, you can view it by running:

```
./draw ../docs/task4.svg
```





screenshot of `svg/basic/test7.svg`

## Task5

### 5.1 What is Pixel Sampling?


During texture mapping, we need to determine the color of each pixel on the screen based on a continuous texture coordinate  $(u, v)$ . Since the texture is stored as a discrete grid of texels, we must "sample" from this grid—this process is called **pixel sampling**. The choice of sampling method directly affects image quality, influencing sharpness, smoothness, and the presence of artifacts such as jaggies or moiré patterns.

### 5.2 Implementation Overview

- In `rasterize_textured_triangle`, barycentric weights are computed and used to interpolate vertex UVs for every subsample, resulting in the target texture coordinate  $(u, v)$ .
- A `SampleParams` structure is filled and passed to the `Texture` at mip level 0.
- The `Texture::sample` function dispatches to either `sample_nearest` or `sample_bilinear` based on the `psm` parameter

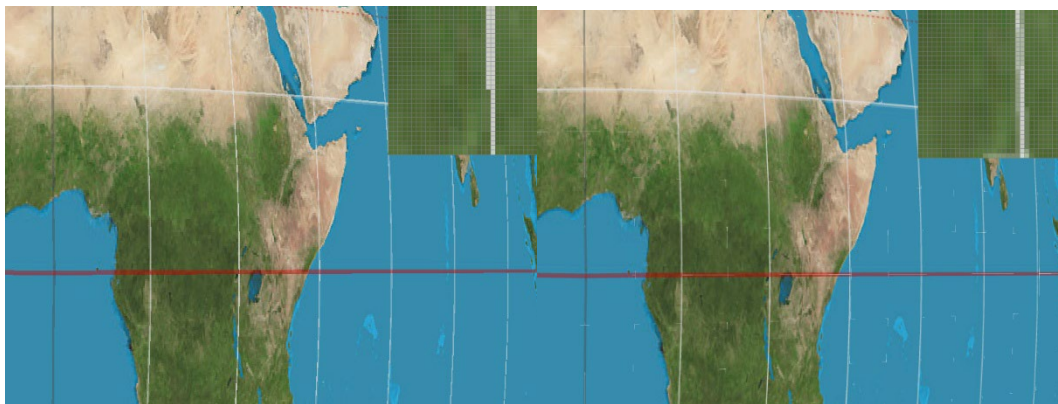
**Nearest:** Converts  $(u, v)$  to texel space, rounds to the nearest integer  $(t_x, t_y)$ , and fetches the single texel using `get_texel`.

**Bilinear:** Fetches the four neighboring texels surrounding  $(x, y)$ , computes the horizontal ( $s_x$ ) and vertical ( $s_y$ ) interpolation weights, and performs two linear interpolations to blend the colors



bilinear sampling at 1 sample

bilinear sampling at 16 sample



nearest sampling at 1 sample

nearest sampling at 16 sample

### 5.3 Discuss the difference

A large difference emerges when a screen pixel spans fractions of a texel—such as during strong magnification of high-frequency textures or thin diagonal edges—because nearest-neighbor snaps to one texel while bilinear interpolates among four. The resulting texel-to-texel discontinuities produce jagged aliasing and moiré in nearest sampling, whereas bilinear’s low-pass blend smooths those high-frequency transitions.

## Task6

### 1 What “level sampling” means

When the same texture is projected onto screen-space, one screen pixel can cover anything from **less than one texel** (magnification) to **hundreds of texels** (strong minification, e.g. a checkerboard far in the distance).



Level sampling decides **which pre-filtered “mipmap” level** of the texture we should read so that the texel density roughly matches the pixel footprint:



## Comparison



Technique you turn up	Speed (per shaded pixel)	Extra memory	Antialiasing ability
<b>Pixel sampling</b>	Nearest = 1 texel fetch Bilinear = 4 fetches	none	Removes blockiness while magnifying, but cannot fight minification aliasing
<b>Level sampling</b>	$L\_NEAREST \approx 1 \times$ pixel cost of chosen $psmL\_LINEAR \approx 2 \times$ pixel cost of psm	All mip levels $\approx 33\%$ extra texture RAM	Eliminates high-frequency flicker when the texture shrinks; tri-linear further hides discontinuities between levels
<b>Samples-per-pixel</b>	$\times sample\_rate$	A larger multi-sampling cache	Kills jagged edges on geometry and picks up true pixel coverage; also improves <i>isotropic</i> texture aliasing, but less efficient for heavy minification

## Screenshot



**L\_NEAREST and P\_LINEAR**

**L\_ZERO and P\_LINEAR**



**L\_ZERO and P\_NEAREST**

**L\_NEAREST and P\_NEAREST**